


Day 5 – Logistic Regression & Classification

 **Date:** 30 June 2025

Today's session was quite interesting as we stepped into the world of **classification problems** in machine learning. Unlike previous regression tasks where we were predicting continuous values, here the focus was on predicting **categories** or **labels**.

Getting Started with Classification

We first understood what classification is — simply put, it's when the output we're trying to predict belongs to specific classes like *yes/no*, *male/female*, or, like in our case today, the **species of flowers**.

A few examples where classification is used:

- Whether an email is spam or not
- Predicting if a patient has a disease
- Classifying images (cat, dog, car, etc.)
- Approving or rejecting a bank loan

Today, we mainly worked with **multi-class classification** using the **Iris dataset**, which involves classifying flowers into one of three species based on four features (like petal length and sepal width).

What We Did Practically

Using Python and the scikit-learn library, we followed these steps:

1. **Loaded the Iris dataset** using `load_iris()`
 2. **Split the data** into training and testing sets
 3. **Built a Logistic Regression model** using `LogisticRegression()`
 4. **Trained the model** and used it to make predictions
 5. **Checked model accuracy** using `accuracy_score()`
 6. Finally, **predicted the species** of a new flower sample
-

 Code We Used (Explained in Simple Steps):

```
python
CopyEdit
from sklearn.datasets import load_iris
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

This part loads all the necessary libraries.

`load_iris()` gives us the dataset, and `LogisticRegression` is our algorithm.

`train_test_split` helps in splitting the data into training and testing sets, while `accuracy_score` helps us check how well the model performed.

python

CopyEdit

```
iris = load_iris()
```

```
X = iris.data # input features
```

```
y = iris.target # output labels (species)
```

```
labels = iris.target_names
```

`X` is the main data (features of flowers), and `y` contains the labels (species).

`labels` is just to get readable names like “setosa”.

python

CopyEdit

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Splits 80% data for training and 20% for testing.

python

CopyEdit

```
model = LogisticRegression(max_iter=200)
```

```
model.fit(X_train, y_train)
```

This trains the model. The `max_iter=200` ensures the algorithm gets enough cycles to find the best solution.

python

CopyEdit

```
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Accuracy: {accuracy * 100:.2f}%")
```

We then predict on the test data and check how accurate the predictions are.

python

CopyEdit

```
sample = [[5.0, 3.6, 1.4, 0.2]]
```

```
prediction = model.predict(sample)
```

```
print(f"Predicted class: {labels[prediction[0]]}")
```

Here, we give the model a new flower with some measurements, and it tells us the species. Pretty cool!



Bonus: Titanic Dataset Introduction

We also had a quick look at the **Titanic dataset** from Kaggle. It's used to predict whether a passenger **survived or not**, based on details like age, gender, ticket class, and fare.

It's a **binary classification** problem, and we'll be using it in upcoming sessions when we learn more algorithms like Decision Trees.



Final Thoughts

This session helped me clearly understand the difference between **classification** and **regression**, and gave hands-on experience with one of the simplest yet powerful classification techniques — logistic regression.