

Day 21: Variance & Gradient Descent – Balancing Accuracy and Learning Speed

Focus of the Day:

We explored two critical concepts today: **Variance**, which affects how well a model generalizes, and **Gradient Descent**, the algorithm powering weight updates in neural networks. Together, they shape how effectively a machine learning model learns and performs on unseen data.

Understanding Variance in Machine Learning

Variance represents **how sensitive a model's predictions are to changes in training data**. It's a crucial part of the **bias-variance tradeoff**, which determines whether your model is underfitting, overfitting, or generalizing well.

Types of Variance

Low Variance (Underfitting):

- Predictions barely change with different datasets.
- Indicates a model that's too simple (e.g., linear model for complex data).
- High bias dominates here, failing to capture patterns.

High Variance (Overfitting):

- Predictions fluctuate wildly even with small data changes.
- The model memorizes noise instead of learning patterns.
- Great on training data, poor on test data.

Optimal Variance (Generalization):

- Balanced bias and variance.
- The model captures essential patterns but ignores noise.
- Achieved via regularization, pruning complexity, or proper hyperparameter tuning.

Analogy:

Imagine shooting arrows:

- Low variance: All arrows land in one spot, but far from the bullseye (systematic error).
 - High variance: Arrows scatter all over, even if some hit near the center.
 - Optimal variance: Arrows cluster tightly around the bullseye.
-

Gradient Descent: The Engine of Learning

Gradient Descent is an optimization algorithm that **drives neural networks to minimize error** by adjusting weights step by step.

How Gradient Descent Works

1. Start with random weights.

2. Calculate the loss (error) for current predictions.
3. Compute gradients (derivatives) of the loss with respect to each weight.
4. Update weights by moving in the **opposite direction of the gradient** (towards lower loss).

Update Rule:

$$w_{\text{new}} = w_{\text{old}} - \alpha \cdot \frac{\partial L}{\partial w}$$

- w : weight
- α : learning rate (step size)
- $\frac{\partial L}{\partial w}$: slope of loss curve w.r.t. weight

Types of Gradient Descent

Batch Gradient Descent:

- Uses the **entire dataset** for one update.
- Stable but slow and memory-heavy.

Stochastic Gradient Descent (SGD):

- Updates weights after **each sample**.
- Faster, but noisy (loss fluctuates).

Mini-Batch Gradient Descent:

- Processes small batches (like 32 or 64 samples).
- Best of both worlds: speed + stability.
- Widely used in deep learning frameworks like TensorFlow and PyTorch.

Connection Between Variance and Gradient Descent

- Poor learning rate \rightarrow model may fail to converge (high bias/variance remains).
- Overly complex networks + aggressive learning rate \rightarrow risk of **high variance (overfitting)**.
- Using **regularization + tuned optimizers (Adam, RMSProp)** can stabilize both variance and convergence speed.

Neural Networks in Action

In deep learning, **backpropagation + gradient descent** iteratively updates weights layer by layer.

- Early layers: learn basic features (edges, colors).
- Deeper layers: capture complex relationships.
- Balanced training avoids both underfitting and overfitting.

Key Takeaways

Variance tells you how stable your model is across different datasets.

The bias-variance tradeoff decides generalization quality.

Gradient Descent is the core algorithm behind all learning in neural networks.